

Operationalizing an AWS ML Project

- Manish Panjwani

Note: For each of the required steps all the snapshots are also present in the **snapshots** folder in the repository.

1. Initial Setup

I have chosen the “**ml.t2.medium**” instance type for Notebook instance (Figure 1 - SageMaker Notebook Instance). There are multiple reasons for selecting this instance type for my notebook.

- Firstly, for completing the execution of this project’s jupyter notebooks we **do not need** a very computationally **powerful CPU** and **high RAM**.
- We will need to keep this **notebook** instance in “**inService**” status for a **long time** while we are working on the project
- In order to **avoid high costs**, we should **select a notebook** that is **low in per hour cost** and **offers** reasonably **good CPU and RAM**.
- So looking at the instance type and their pricing: <https://aws.amazon.com/SageMaker/pricing/>, we have two choices “ml.t2.medium” and “ml.t3.medium”. Both have 2 vCPU and 4 GB Memory, and as per doc “ml.t3.medium” has a slightly higher cost as it has a fast boot time.
- Now given that we do have a critical requirement for a fast boot time, so we can go ahead with the “**ml.t2.medium**” as it offers the same computational power and is lower in per hour cost.




Figure 1. SageMaker Notebook Instance

The dog breed dataset was uploaded to a newly created S3 bucket, successfully.

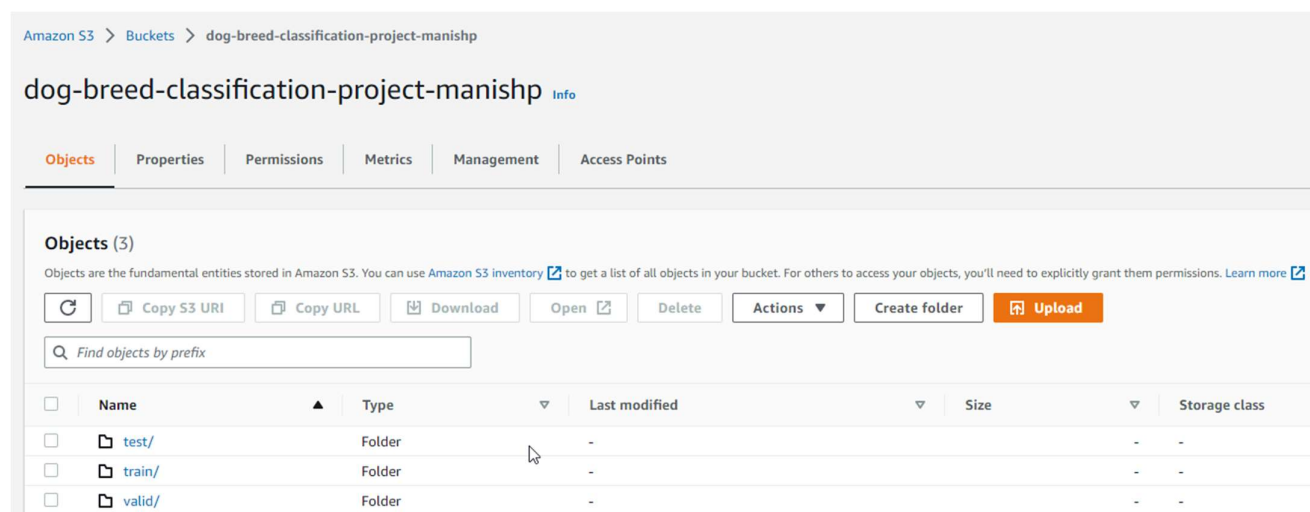


Figure 2. S3 Bucket snapshot

2. SageMaker Training and Deployment

For hyperparameter tuning I tried to increase the processing power a bit by using the “ml.m5.2xlarge” for the single instance and multi-instance training purposes.

upon training the model with the best parameters from above tuning, the model gave a 0 test accuracy! So increased the max_jobs = 6, max_parallel_jobs = 3 and also changed its instance_type = “ml.m5.xlarge” to speed up the computations a bit.

Ran the hyperparameter jobs and it executed successfully:

Name	Status	Training completed/total	Creation time	Duration
pytorch-training-230102-2321	Completed	6 / 6	Jan 02, 2023 23:21 UTC	27 minutes

Figure 3. Hyperparameter jobs summary

Post which triggered the **single instance** and **multi-instance training jobs**. Jobs completed successfully. For snapshot of the training jobs refer the images folder in the repository.

Job name dog-pytorch-2023-01-03-00-37-33-365	Status Completed View history	SageMaker metrics time series Enabled	IAM role ARN arn:aws:iam::412977242314:role/service-role/AmazonSageMaker-ExecutionRole-20220807T073175
ARN arn:aws:sagemaker:us-east-1:412977242314:training-job/dog-pytorch-2023-01-03-00-37-33-365	Creation time Jan 03, 2023 00:37 UTC	Training time (seconds) 695	
	Last modified time Jan 03, 2023 00:50 UTC	Billable time (seconds) 695	
		Managed spot training savings 0%	
		Tuning job source/parent -	

Algorithm			
Algorithm ARN -	Additional volume size (GB) 30	Maximum wait time for managed spot training(s) -	Volume encryption key -
Training image 763104351884.dkr.ecr.us-east-1.amazonaws.com/pytorch-training:1.4.0-cpu-py3	Maximum runtime (s) 86400	Managed spot training Disabled	
Input mode File			

Instance group	Instance type	Instance count	Keep alive period
-	ml.m5.2xlarge	1	-

Figure 4. Single Instance Training Job

Job name dog-pytorch-2023-01-03-01-02-31-722	Status ✔ Completed View history	SageMaker metrics time series Enabled	IAM role ARN arn:aws:iam::412977242314:role/AmazonSageMaker-Exec-20220807T073175
ARN arn:aws:sagemaker:us-east-1:412977242314:training-job/dog-pytorch-2023-01-03-01-02-31-722	Creation time Jan 03, 2023 01:02 UTC	Training time (seconds) 722	
	Last modified time Jan 03, 2023 01:17 UTC	Billable time (seconds) 722	
		Managed spot training savings 0%	
		Tuning job source/parent -	

Algorithm			
Algorithm ARN -	Additional volume size (GB) 30	Maximum wait time for managed spot training(s) -	Volume encryption key -
Training image 763104351884.dkr.ecr.us-east-1.amazonaws.com/pytorch-training:1.4.0-cpu-py3	Maximum runtime (s) 86400	Managed spot training Disabled	
Input mode File			

Instance group	Instance type	Instance count	Keep alive period
-	ml.m5.2xlarge	5	-

Figure 5. Multi Instance Training Job

Deployed Endpoints:

- Single instance deployed endpoint: “pytorch-inference-2023-01-03-00-56-36-940”
- Multi instance deployed endpoint: “pytorch-inference-2023-01-03-01-25-08-339”

Amazon SageMaker > Endpoints					
Endpoints		Refresh	Update endpoint	Actions	Create endpoint
<input type="text" value="Search endpoints"/>		< 1 > Settings			
	Name	ARN	Creation time	Status	Last updated
<input type="radio"/>	pytorch-inference-2023-01-03-01-25-08-339	arn:aws:sagemaker:us-east-1:412977242314:endpoint/pytorch-inference-2023-01-03-01-25-08-339	Jan 03, 2023 01:25 UTC	✔ InService	Jan 03, 2023 01:27 UTC
<input type="radio"/>	pytorch-inference-2023-01-03-00-56-36-940	arn:aws:sagemaker:us-east-1:412977242314:endpoint/pytorch-inference-2023-01-03-00-56-36-940	Jan 03, 2023 00:56 UTC	✔ InService	Jan 03, 2023 00:58 UTC

Figure 6. SageMaker Endpoints

3. EC2 Training

- We have utilized the **t2.xlarge** instance and the **Deep Learning AMI GPU PyTorch 1.13.1 (Amazon Linux 2)**. This seems like a reasonable balance of performance and affordability.
- As per the documentation, T2 instances can sustain high CPU performance for as long as a workload needs it.
- For most general-purpose workloads, T2 instances will provide ample performance without any additional charges.
- Similarly, because we don't know the duration for which we might need to keep this EC2 instance running for training, it's better to go with a medium size instance so we don't have to pay for a large instance while we're doing setup, debugging and other tasks.
- I have to install missing modules like numpy, torch, torchvision, tqdm

Difference between ec2train1.py (EC2 script) and train_and_deploy-solution.ipynb + hpo.py (SageMaker scripts)


```

[root@ip-172-31-26-183 TrainedModels]# ls -lrt
total 93212
-rw-r--r-- 1 root root 95445365 Jan  3 04:49 model.pth
[root@ip-172-31-26-183 TrainedModels]# pwd
/root/TrainedModels
[root@ip-172-31-26-183 TrainedModels]# echo `hostname`
ip-172-31-26-183.ec2.internal
[root@ip-172-31-26-183 TrainedModels]#

```

Figure 10. EC2 Training saved - model.pth

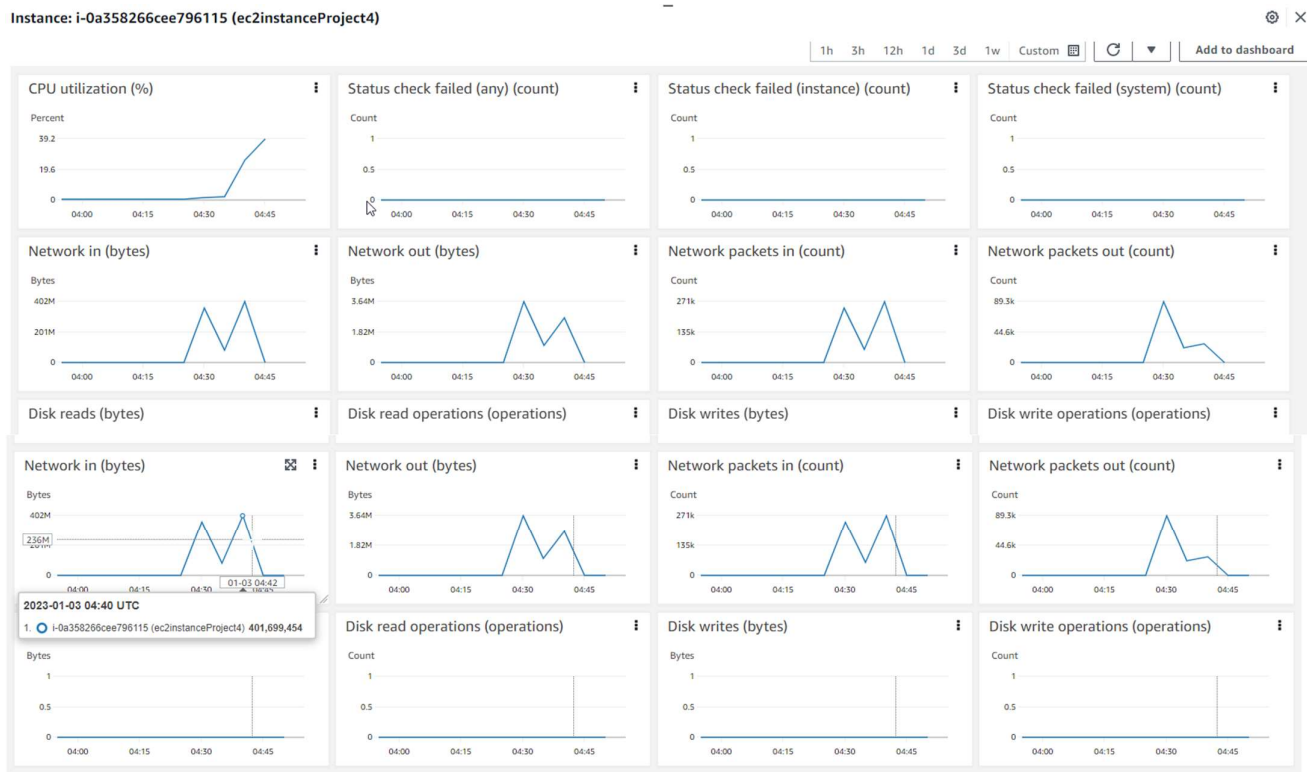


Figure 11. EC2 Instance Resource Metrics

4. Lambda functions

- The lambda functions will be used for invoking our deployed endpoints.
- The lambda function implemented in this project expects the image inputs in json format, which is used to invoke the model's deployed endpoint
- Given we have two endpoints deployed, one for the single instance training and the other for the multi-instance training, we will only use the multi-instance training jobs endpoint and create a lambda function for invoking that endpoint.
- Multi instance trained endpoint that we will be using: **"pytorch-inference-2023-01-03-01-25-08-339"**
- We created the lambda function with the corresponding changes to invoke the endpoint:




Figure 12. Lambda Function

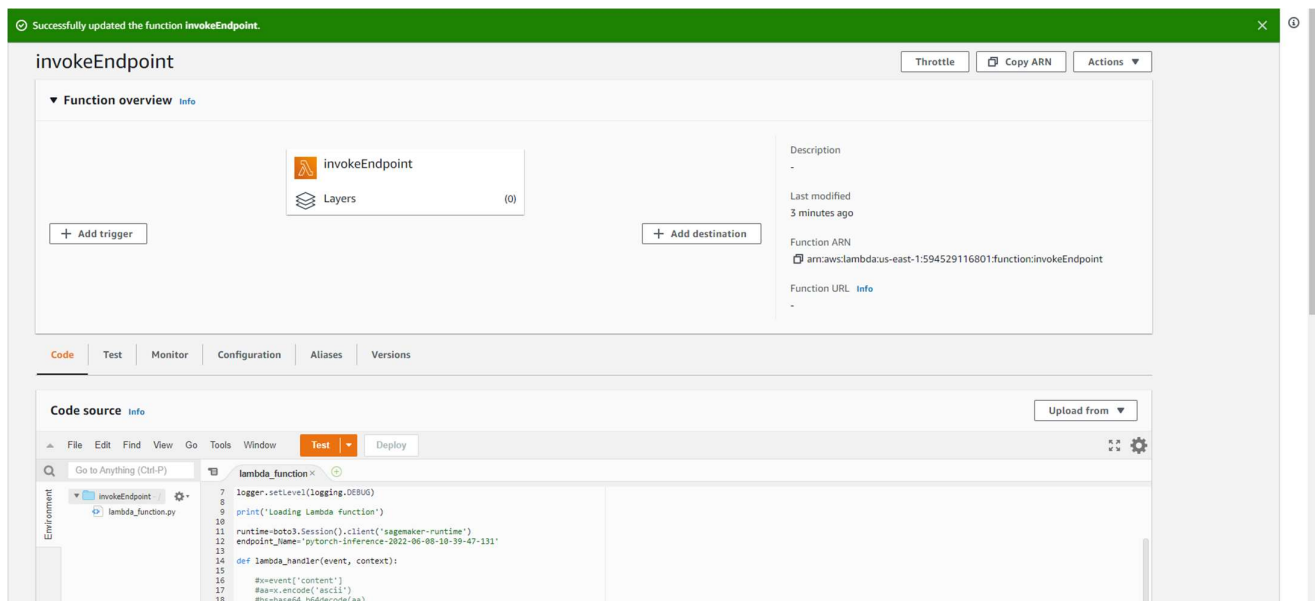


Figure 13. Lambda Function

5. Security and Testing

- We had created a new role for this lambda functions with basic accesses.
- We used the below test event to test our lambda function:

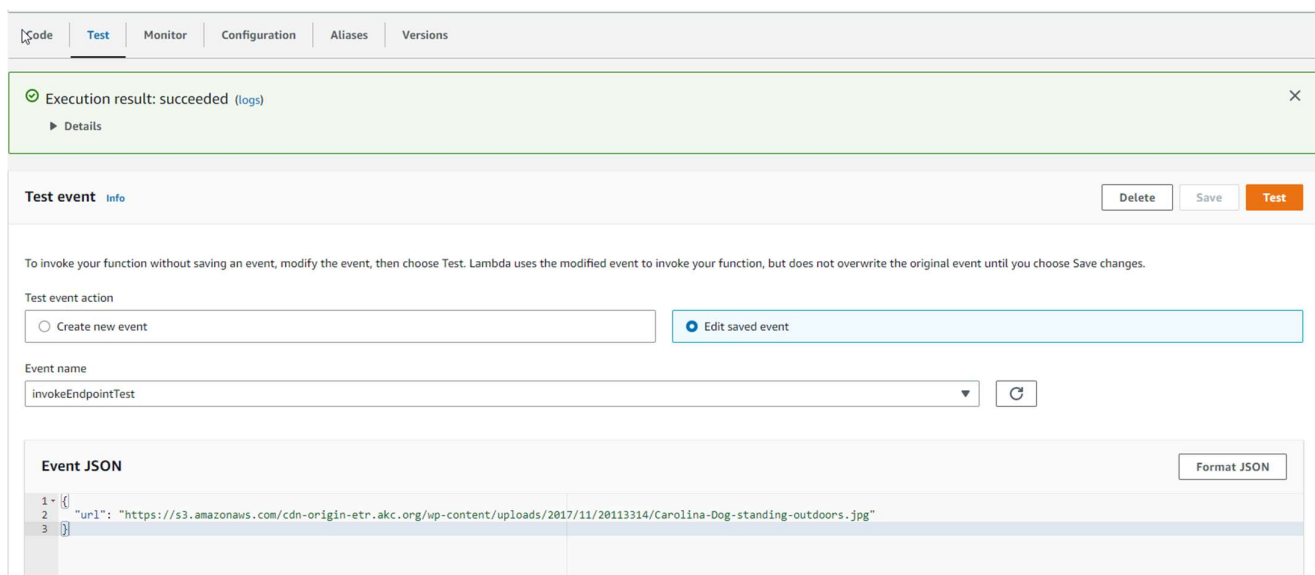


Figure 14. Lambda Function Test Event

- Now when we tried to execute the test event we got and **AccessDeniedException**. This was expected as the lambda function did not have access to invoke the SageMaker endpoint.
- Error Message:

```
{
  "errorMessage": is not authorized to perform:
}
```

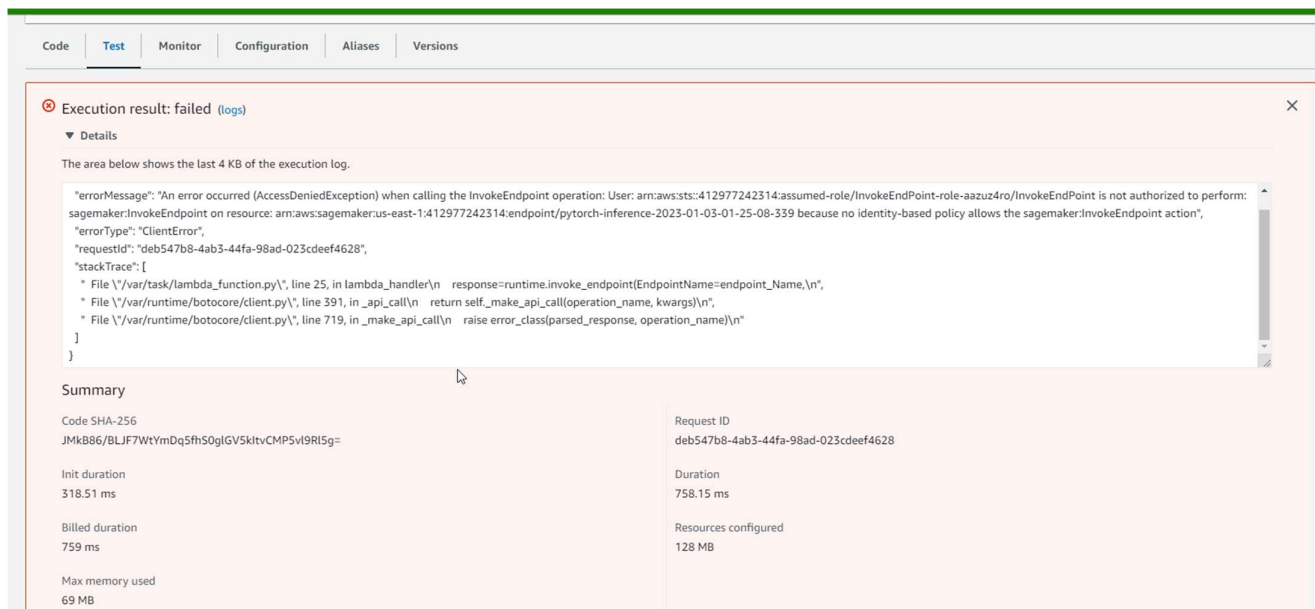



Figure 15. Lambda function test event failure response

- So went ahead and added the “SageMakerFullAccess” policy role to the lambda function’s role.

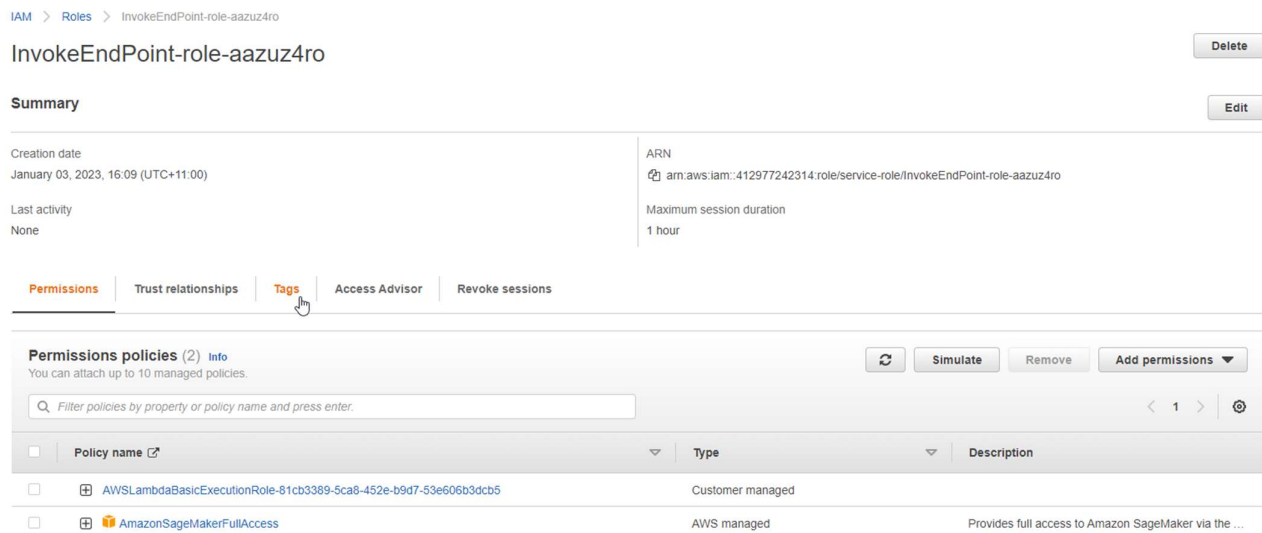


Figure 16. Lambda function role IAM permissions

- Post which we got the following output from the test event:
- (Please note that there are 133 dog breed and not 33 as mentioned in the project instructions. So, we do expect there to be around 133 elements in the response object.)
- Response:

```
{
  "statusCode": 200,
  "headers": {
    "Content-Type": "text/plain",
    "Access-Control-Allow-Origin": "*"
  },
  "type-result": "<class 'str'>",
  "Content-Type-In": "LambdaContext ([aws_request_id=e6839e8c-899e-4c74-a871-3bcla8f19302, log_group_name=/aws/lambda/InvokeEndPoint, log_stream_name=2023/01/03/[$LATEST]e1124ef245b640389dd9ae2feld459dc, function_name=InvokeEndPoint, memory_limit_in_mb=128, function_version=$LATEST, invoked_function_arn=arn:aws:lambda:us-east-1:412977242314:function:InvokeEndPoint, client_context=None, identity=CognitoIdentity ([cognito_identity_id=None, cognito_identity_pool_id=None]))",
}
```

```

"body": "[[-8.98412036895752, -3.4478538036346436, -3.629460096359253, -
1.2364776134490967, -2.606005907058716, -3.641772747039795, -4.87917947769165, -
2.765993356704712, -5.374297142028809, -0.2598344683647156, 0.36318743228912354, -
3.9999375343322754, -3.839862108230591, 0.3446093797683716, -9.077698707580566, -
5.279119491577148, -9.107315063476562, -3.7618331909179688, -3.669107675552368, -
1.516728162765503, -5.2369890213012695, -2.1203291416168213, -7.924822807312012, -
4.405004501342773, -5.5941267013549805, -7.448276042938232, -4.458913803100586, -
4.156890392303467, -6.0630974769592285, -1.9943898916244507, -2.9814293384552, -
3.0918822288513184, -8.973814964294434, -4.173212051391602, -7.15018892288208, -
8.212821960449219, -5.752190113067627, -4.245693683624268, -1.4357178211212158, -
4.766050338745117, -3.1229248046875, -3.886951446533203, -0.6014687418937683, -
4.205016613006592, 0.34389206767082214, -8.061474800109863, -1.937998652458191, -
0.24757209420204163, -2.6457574367523193, -1.7174931764602661, -7.95383882522583, -
5.795444011688232, -5.12360954284668, -3.567878246307373, -7.9134111404418945, -
1.7634156942367554, -8.026968002319336, -6.81389045715332, -3.278583526611328, -
3.9052999019622803, -6.31833553314209, -7.021487712860107, -10.082239151000977, -
7.719013214111328, -6.309490203857422, -7.944791793823242, -0.09204558283090591, -
5.567389965057373, -1.8296440839767456, -1.0008320808410645, -0.014354053884744644, -
6.184256076812744, -6.5170979499816895, -7.322001934051514, -6.667438507080078, -
4.668673038482666, -8.859840393066406, -4.342897891998291, -5.429651737213135, -
6.887331962585449, -1.463620901107788, -10.355123519897461, -2.3695499897003174, -
0.6696080565452576, -8.244217872619629, -3.960181713104248, -4.961513519287109, -
6.852430820465088, -4.977674961090088, -1.9244245290756226, -6.798636436462402, -
4.3571553230285645, -7.417713165283203, -7.065046310424805, -3.170226812362671, -
2.61667799949646, -8.007290840148926, -5.260977745056152, -7.49748420715332, -
5.690293788909912, -8.133918762207031, -1.8966071605682373, -1.5495470762252808, -
7.87912654876709, -6.289516925811768, -6.853586196899414, -3.980215549468994, -
0.08247441053390503, -2.2983858585357666, -2.2757515907287598, -2.413423538208008, -
2.8458051681518555, -8.328771591186523, -6.05941104888916, -4.543338775634766, -
2.134174108505249, -6.622488498687744, -1.2265808582305908, -9.922281265258789, -
1.0559980869293213, -3.152406692504883, -2.850202798843384, -4.3935441970825195, -
2.4728634357452393, -8.680174827575684, -4.717899799346924, -4.835120677947998, -
1.554640293121338, -5.1974029541015625, -9.266709327697754, -9.486908912658691, -
2.4802603721618652, -5.181535720825195]]]"
}

```

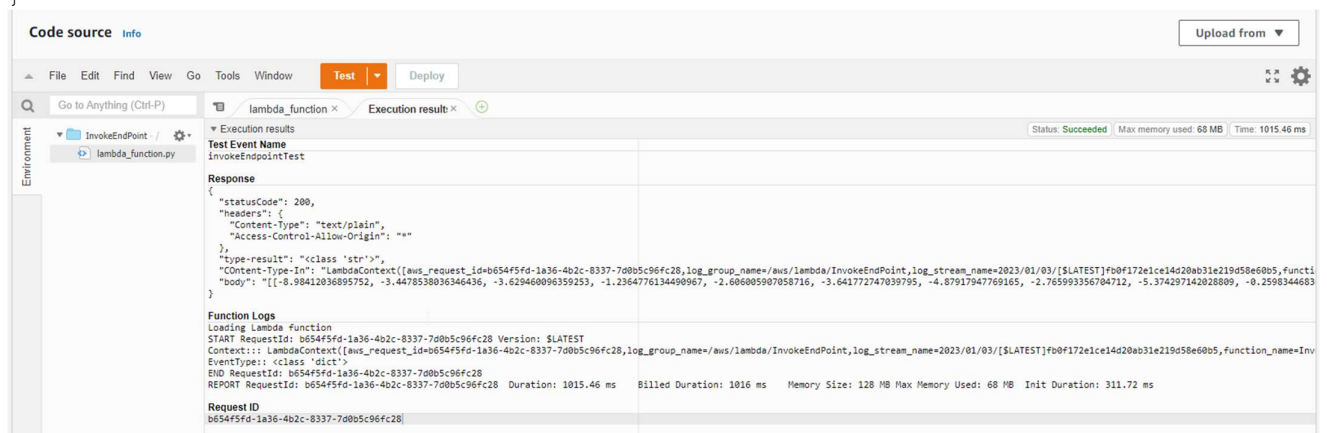


Figure 17. Lambda function test event success response

- I'm concerned about the "Full Access" type permission policies that are available.
- For example, for this lambda functions we have provided the lambda function a Full Access to SageMaker resources, but this does not seem to follow the concept of least privilege access.
- Ideally, one should only allow these lambda functions to query the endpoints that they're intended and allowed to query.
- We will have to do some more analysis to figure out if there's anything we could do about it.
- Furthermore, another concern is that the account's root user does not employ MFA
- Looking at the IAM roles that are currently active, all the roles seem to be necessary and most of the roles have been added on a per need basis.

- However, we need to keep an eye on the roles dashboard to make sure only relevant roles necessary for currently active projects, are the only roles that are in active state to prevent unauthorized accesses.

The screenshot shows the AWS IAM dashboard. On the left is a navigation menu with options like 'Dashboard', 'Access management', 'Users', 'Roles', 'Policies', 'Identity providers', 'Account settings', 'Access reports', 'Access analyzer', 'Archive rules', 'Analyzers', 'Settings', and 'Credential report'. The main content area is titled 'IAM dashboard'. It features a 'Security recommendations' section with a red alert icon and a '1' badge, indicating one recommendation. Below this is a 'What's new' section with updates for features in IAM, including advanced notices about Amazon S3, AWS IAM Identity Center, AWS Lambda, and Amazon ElastiCache. The 'IAM resources' section displays a summary of IAM entities: 0 User groups, 0 Users, 26 Roles, 18 Policies, and 0 Identity providers.

Figure 18. IAM Dashboard

The screenshot shows the AWS IAM Roles page. The left navigation menu is the same as in Figure 18. The main content area is titled 'Roles (26) Info'. It includes a search bar and a table of roles. The table has columns for 'Role name', 'Trusted entities', and 'Last activity'. The roles listed include 'AmazonSageMaker-ExecutionRole-20220807T073175', 'AWSServiceRoleForAmazonSageMakerNotebooks', 'AWSServiceRoleForAWSCloud9', 'AWSServiceRoleForCloudWatchEvents', 'AWSServiceRoleForElasticCache', 'AWSServiceRoleForOrganizations', 'AWSServiceRoleForSupport', 'AWSServiceRoleForTrustedAdvisor', 'EMR_AutoScaling_DefaultRole', 'EMR_DefaultRole', 'EMR_EC2_DefaultRole', and 'ImageClassifier-role-6xo0eur0'. Each role has a checkbox for selection and a link to its details page.

Figure 19. IAM Roles

6. Concurrency and Auto-scaling

- Before adding in configs for Concurrency and Auto-scaling for our lambda functions we will first create a version config for our lambda function.

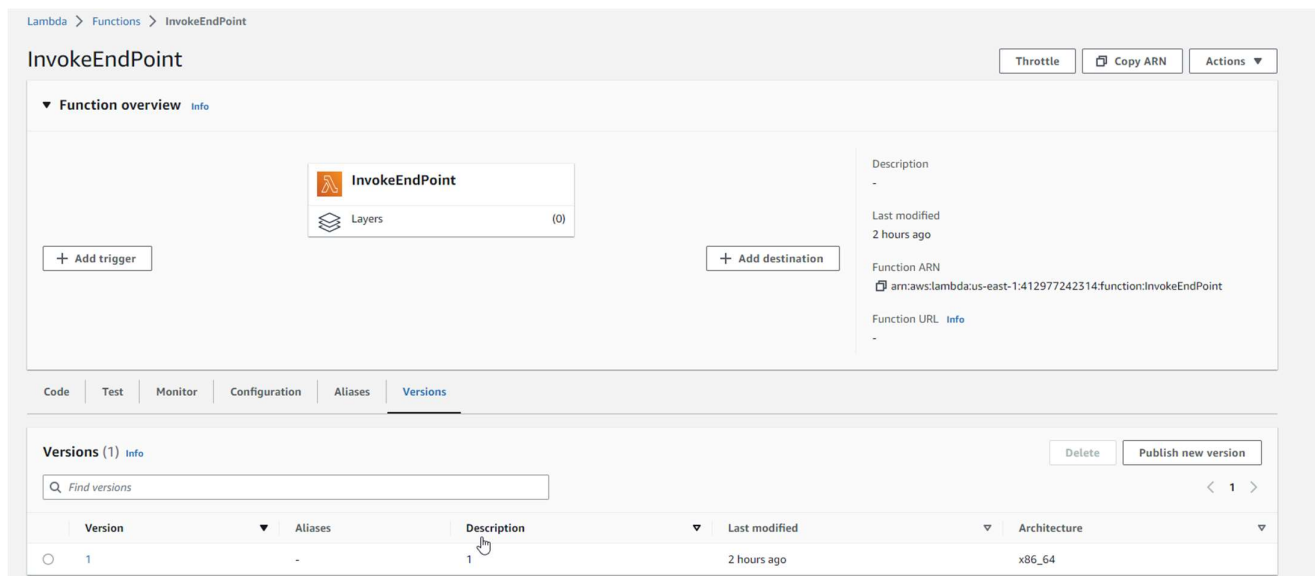


Figure 20. Lambda function version config

- For the lambda function we have set the **reserved concurrency** to be 5. This implies that the lambda function would be able to handle up to 5 requests concurrently at the same time. This would help lower latency issues in situations when there is higher traffic than usual.

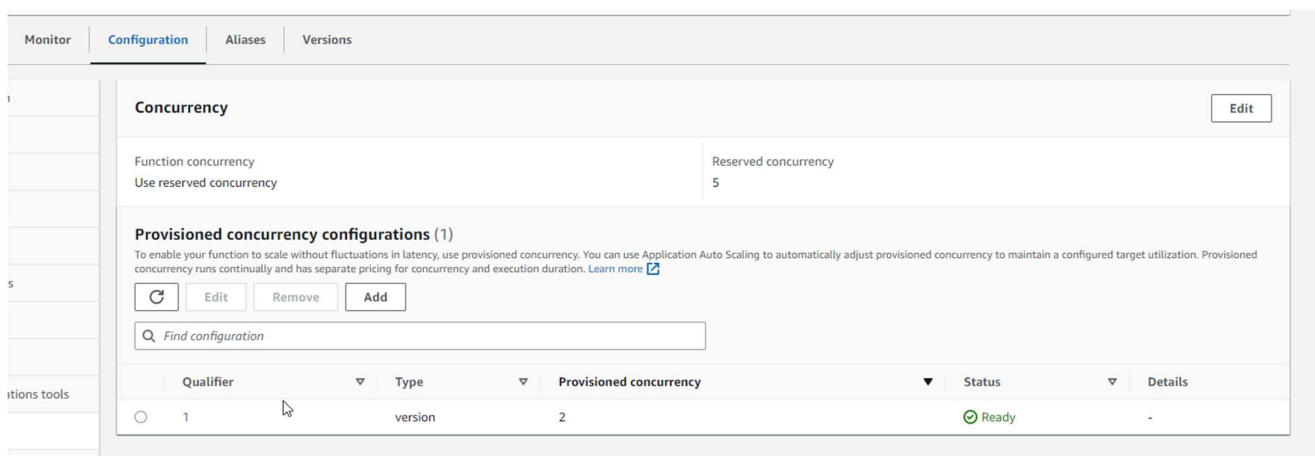


Figure 21. lambda function reversed concurrency

- Given the current use case, ideally using only reserved concurrency should have sufficed for our use case and we might not need to consider using the provisioned concurrency configs. However, for the sake of completion, I tried to add in the config for the provisioned concurrency as well. We set the provision concurrency to be 2.

Concurrency Edit

Function concurrency
Use reserved concurrency

Reserved concurrency
5

Provisioned concurrency configurations (1)
To enable your function to scale without fluctuations in latency, use provisioned concurrency. You can use Application Auto Scaling to automatically adjust provisioned concurrency to maintain a configured target utilization. Provisioned concurrency runs continually and has separate pricing for concurrency and execution duration. [Learn more](#)

Refresh Edit Remove Add

	Qualifier	Type	Provisioned concurrency	Status	Details
<input type="radio"/>	1	version	2	Ready	-

Figure 22. lambda function provisioned concurrency

- For adding config for auto-scaling we have added the below configs:

Minimum instance count: 1 - Maximum instance count: 3

IAM role
Amazon SageMaker uses the following service-linked role for automatic scaling. [Learn more](#)
AWSServiceRoleForApplicationAutoScaling_SageMakerEndpoint

Built-in scaling policy [Learn more](#)

Policy name
SageMakerEndpointInvocationScalingPolicy

Target metric: SageMakerVariantInvocationsPerInstance [Learn more](#) Target value: 10

Scale in cool down (seconds) - optional: 30 Scale out cool down (seconds) - optional: 30

☐ Disable scale in
Select if you don't want automatic scaling to delete instances when traffic decreases. [Learn more](#)

Figure 23. Endpoint Auto-scaling Config

- We have set the max instance count to 3 for Auto-scaling, as considering the current requirement, auto scaling on 3 instances with a scale-in and scale-out cool down time of 30 seconds should be a reasonably good config.

Endpoint runtime settings

Update weights

Update instance count

Configure auto scaling

		Variant name ▲	Current weight ▼	Desired weight	Elastic Inference	Instance type ▼	Current instance count ▼	Desired instance count ▼	Instance min - max	Automatic scaling
<input type="radio"/>	<input checked="" type="radio"/>	AllTraffic	1	1	-	ml.m5.large	1	1	1 - 3	Yes

Figure 24. Endpoint Auto-scaling Metrics